

# Fastag-fraud-detection

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
from sklearn.metrics import classification_report, accuracy_score, precision_score,
from imblearn.over_sampling import SMOTE
```

## 1. Project Overview


This project focuses on developing a robust machine learning- based fraud detection system for fastag transactions. Fastag is an electronic toll collection system in India that uses RFID technology to make toll payments directly from a prepaid account linked to a user's vehicle. As digital transactions become machine learning classification techniques to accurately also increases. This Project aims to leverage machine learning classification techniques to accurately identify fraudulent transactions, thereby ensuring the security and integrity of fastag Transactions.

```
In [2]: df = pd.read_csv(r'C:\Users\chira\Downloads\FastagFraudDetection.csv')
df
```

Out[2]:

	Transaction_ID	Timestamp	Vehicle_Type	FastagID	TollBoothID	Lane_Type	Vehicle_Dimensi
0	1	01-06-2023 11:20	Bus	FTG-001- ABC-121	A-101	Express	Li
1	2	01-07-2023 14:55	Car	FTG-002- XYZ-451	B-102	Regular	S
2	3	01-08-2023 18:25	Motorcycle	NaN	D-104	Regular	S
3	4	01-09-2023 02:05	Truck	FTG-044- LMN- 322	C-103	Regular	Li
4	5	01-10-2023 06:35	Van	FTG-505- DEF-652	B-102	Express	Mec
...	...	...	...	...	...	...	...
4995	4996	01-01-2023 22:18	Truck	FTG-445- EDC-765	C-103	Regular	Li
4996	4997	1/17/2023 13:43	Van	FTG-446- LMK-432	B-102	Express	Mec
4997	4998	02-05-2023 05:08	Sedan	FTG-447- PLN-109	A-101	Regular	Mec
4998	4999	2/20/2023 20:34	SUV	FTG-458- VFR-876	B-102	Express	Li
4999	5000	03-10-2023 00:59	Bus	FTG-459- WSX- 543	C-103	Regular	Li

5000 rows × 13 columns



## 2.Dataset Description

The Dataset Comprises Various Features related to fastag Transcation, including transactiondetails, vehicle information, geographical location, and transcation amounts. The key features are:

Transcation\_ID : Unique Identifier fro each transaction.

Timestamp: Date and time of the transaction.

Vehicle\_Type: Type of vehicle involved in the transaction.

FastagID: Unique identifier for Fastag.

TollBoothID: Identifier for the toll booth.

Lane\_Type: Type of lane used for the transaction.

Vehicle\_Dimensions: Dimensions of the vehicle.

Transaction\_Amount: Amount associated with the transaction.

Amount\_paid: Amount paid for the transaction.

Geographical\_Location: Location details of the transaction.

Vehicle\_Speed: Speed of the vehicle during the transaction.

Vehicle\_Plate\_Number: License plate number of the vehicle.

Fraud\_indicator: Binary indicator of fraudulent activity (target variable).

Timestamp: Date and time of the transaction.

Vehicle\_Type: Type of vehicle involved in the transaction.

FastagID: Unique identifier for Fastag.

TollBoothID: Identifier for the toll booth.

Lane\_Type: Type of lane used for the transaction.

Vehicle\_Dimensions: Dimensions of the vehicle.

Transaction\_Amount: Amount associated with the transaction.

Amount\_paid: Amount paid for the transaction.

Geographical\_Location: Location details of the transaction.

Vehicle\_Speed: Speed of the vehicle during the transaction.

Vehicle\_Plate\_Number: License plate number of the vehicle.

Fraud\_indicator: Binary indicator of fraudulent activity (target variable).

In [3]: `df.head()`

Out[3]:

	Transaction_ID	Timestamp	Vehicle_Type	FastagID	TollBoothID	Lane_Type	Vehicle_Dimensions
0	1	01-06-2023 11:20	Bus	FTG-001- ABC-121	A-101	Express	Large
1	2	01-07-2023 14:55	Car	FTG-002- XYZ-451	B-102	Regular	Small
2	3	01-08-2023 18:25	Motorcycle	NaN	D-104	Regular	Small
3	4	01-09-2023 02:05	Truck	FTG-044- LMN-322	C-103	Regular	Large
4	5	01-10-2023 06:35	Van	FTG-505- DEF-652	B-102	Express	Medium

In [4]: `df.tail()`

Out[4]:	Transaction_ID	Timestamp	Vehicle_Type	FastagID	TollBoothID	Lane_Type	Vehicle_Dimensi
<b>4995</b>	4996	01-01-2023 22:18	Truck	FTG-445- EDC-765	C-103	Regular	Li
<b>4996</b>	4997	1/17/2023 13:43	Van	FTG-446- LMK-432	B-102	Express	Mec
<b>4997</b>	4998	02-05-2023 05:08	Sedan	FTG-447- PLN-109	A-101	Regular	Mec
<b>4998</b>	4999	2/20/2023 20:34	SUV	FTG-458- VFR-876	B-102	Express	Li
<b>4999</b>	5000	03-10-2023 00:59	Bus	FTG-459- WSX- 543	C-103	Regular	Li

In [5]: `df.sample(5)`

Out[5]:	Transaction_ID	Timestamp	Vehicle_Type	FastagID	TollBoothID	Lane_Type	Vehicle_Dimensi
<b>4119</b>	4120	3/29/2023 14:30	Motorcycle	NaN	D-106	Regular	S
<b>2979</b>	2980	11/22/2023 0:58	Truck	FTG-480- TGB-250	C-103	Regular	Li
<b>642</b>	643	04-08-2023 13:55	Sedan	FTG-823- NMK- 365	A-101	Express	Mec
<b>2644</b>	2645	02-04-2023 06:25	Van	FTG-125- DCF-765	B-102	Express	Mec
<b>2036</b>	2037	9/13/2023 4:05	SUV	FTG-455- QRS-789	B-102	Express	Li

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Transaction_ID                        5000 non-null   int64
1   Timestamp                            5000 non-null   object
2   Vehicle_Type                         5000 non-null   object
3   FastagID                             4451 non-null   object
4   TollBoothID                          5000 non-null   object
5   Lane_Type                            5000 non-null   object
6   Vehicle_Dimensions                   5000 non-null   object
7   Transaction_Amount                   5000 non-null   int64
8   Amount_paid                          5000 non-null   int64
9   Geographical_Location                 5000 non-null   object
10  Vehicle_Speed                         5000 non-null   int64
11  Vehicle_Plate_Number                 5000 non-null   object
12  Fraud_indicator                       5000 non-null   object
dtypes: int64(4), object(9)
memory usage: 507.9+ KB
```

# Summary Statistics

We use summary statistics to get an overview of the numerical features.

## Numerical Summary Statistics:

Count: The Number Of Non-missing Values.

Mean: The Average value.

Std: The Standard Deviation, indicating the spread of the values.

min: The Minimum Value.

25%: The 25th percentile value (first quartile).

50%: The median value(second quartile).

75%: The 75th percentile value (third quartile).

max: The maximum value.

median: The median value, explicitly added for clarity.

mode: The most frequently occurring value.

missing\_values: The count of missing values in each column.

```
In [7]: df.describe()
```

	Transaction_ID	Transaction_Amount	Amount_paid	Vehicle_Speed
count	5000.000000	5000.000000	5000.000000	5000.000000
mean	2500.500000	161.06200	141.261000	67.851200
std	1443.520003	112.44995	106.480996	16.597547
min	1.000000	0.000000	0.000000	10.000000
25%	1250.750000	100.00000	90.000000	54.000000
50%	2500.500000	130.00000	120.000000	67.000000
75%	3750.250000	290.00000	160.000000	82.000000
max	5000.000000	350.00000	350.000000	118.000000

```
In [8]: df.select_dtypes("number").mean()
```

```
Out[8]: Transaction_ID      2500.5000
Transaction_Amount      161.0620
Amount_paid             141.2610
Vehicle_Speed           67.8512
dtype: float64
```

```
In [9]: df.select_dtypes("number").median()
```

```
Out[9]: Transaction_ID      2500.5  
Transaction_Amount      130.0  
Amount_paid            120.0  
Vehicle_Speed          67.0  
dtype: float64
```

```
In [10]: df.select_dtypes('number').mode().iloc[0]
```

```
Out[10]: Transaction_ID      1.0  
Transaction_Amount      0.0  
Amount_paid            0.0  
Vehicle_Speed          55.0  
Name: 0, dtype: float64
```

```
In [11]: df.isnull().sum()
```

```
Out[11]: Transaction_ID      0  
Timestamp      0  
Vehicle_Type    0  
FastagID       549  
TollBoothID     0  
Lane_Type      0  
Vehicle_Dimensions  0  
Transaction_Amount  0  
Amount_paid     0  
Geographical_Location  0  
Vehicle_Speed    0  
Vehicle_Plate_Number  0  
Fraud_indicator  0  
dtype: int64
```

```
In [12]: df["FastagID"].fillna(df["FastagID"].mode()[0], inplace=True)
```

```
In [13]: df.isnull().sum()
```

```
Out[13]: Transaction_ID      0  
Timestamp      0  
Vehicle_Type    0  
FastagID       0  
TollBoothID     0  
Lane_Type      0  
Vehicle_Dimensions  0  
Transaction_Amount  0  
Amount_paid     0  
Geographical_Location  0  
Vehicle_Speed    0  
Vehicle_Plate_Number  0  
Fraud_indicator  0  
dtype: int64
```

```
In [14]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5000 entries, 0 to 4999
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                -
0   Transaction_ID                        5000 non-null   int64
1   Timestamp                            5000 non-null   object
2   Vehicle_Type                         5000 non-null   object
3   FastagID                            5000 non-null   object
4   TollBoothID                         5000 non-null   object
5   Lane_Type                           5000 non-null   object
6   Vehicle_Dimensions                  5000 non-null   object
7   Transaction_Amount                  5000 non-null   int64
8   Amount_paid                         5000 non-null   int64
9   Geographical_Location                5000 non-null   object
10  Vehicle_Speed                       5000 non-null   int64
11  Vehicle_Plate_Number                5000 non-null   object
12  Fraud_indicator                     5000 non-null   object
dtypes: int64(4), object(9)
memory usage: 507.9+ KB

```

```
In [15]: df.drop_duplicates(inplace = True)
df
```

```
Out[15]:
```

	Transaction_ID	Timestamp	Vehicle_Type	FastagID	TollBoothID	Lane_Type	Vehicle_Dimensi
0	1	01-06-2023 11:20	Bus	FTG-001- ABC-121	A-101	Express	Li
1	2	01-07-2023 14:55	Car	FTG-002- XYZ-451	B-102	Regular	S
2	3	01-08-2023 18:25	Motorcycle	FTG-000- QAZ-210	D-104	Regular	S
3	4	01-09-2023 02:05	Truck	FTG-044- LMN- 322	C-103	Regular	Li
4	5	01-10-2023 06:35	Van	FTG-505- DEF-652	B-102	Express	Mec
...	...	...	...	...	...	...	...
4995	4996	01-01-2023 22:18	Truck	FTG-445- EDC-765	C-103	Regular	Li
4996	4997	1/17/2023 13:43	Van	FTG-446- LMK-432	B-102	Express	Mec
4997	4998	02-05-2023 05:08	Sedan	FTG-447- PLN-109	A-101	Regular	Mec
4998	4999	2/20/2023 20:34	SUV	FTG-458- VFR-876	B-102	Express	Li
4999	5000	03-10-2023 00:59	Bus	FTG-459- WSX- 543	C-103	Regular	Li


5000 rows × 13 columns

## 5. Data Visualization

## 6. Visualize the distribution of the target variable

```
In [16]: plt.figure(figsize= (8,6))
sns.countplot(df['Fraud_indicator'])
plt.title('Distribution of fraud Indicator')
plt.xlabel("fraud Indicator")
plt.ylabel("count")
plt.show
```

```
Out[16]: <function matplotlib.pyplot.show(close=None, block=None)>
```




## 7. Visualize the Distribution of Transaction Amounts

```
In [17]: plt.figure(figsize=(10,6))
sns.histplot(df['Transaction_Amount'], bins=30 ,kde = True)
plt.title("distribution of Transaction Amounts")
plt.xlabel("transaction amounts")
plt.ylabel("count")
plt.show
```

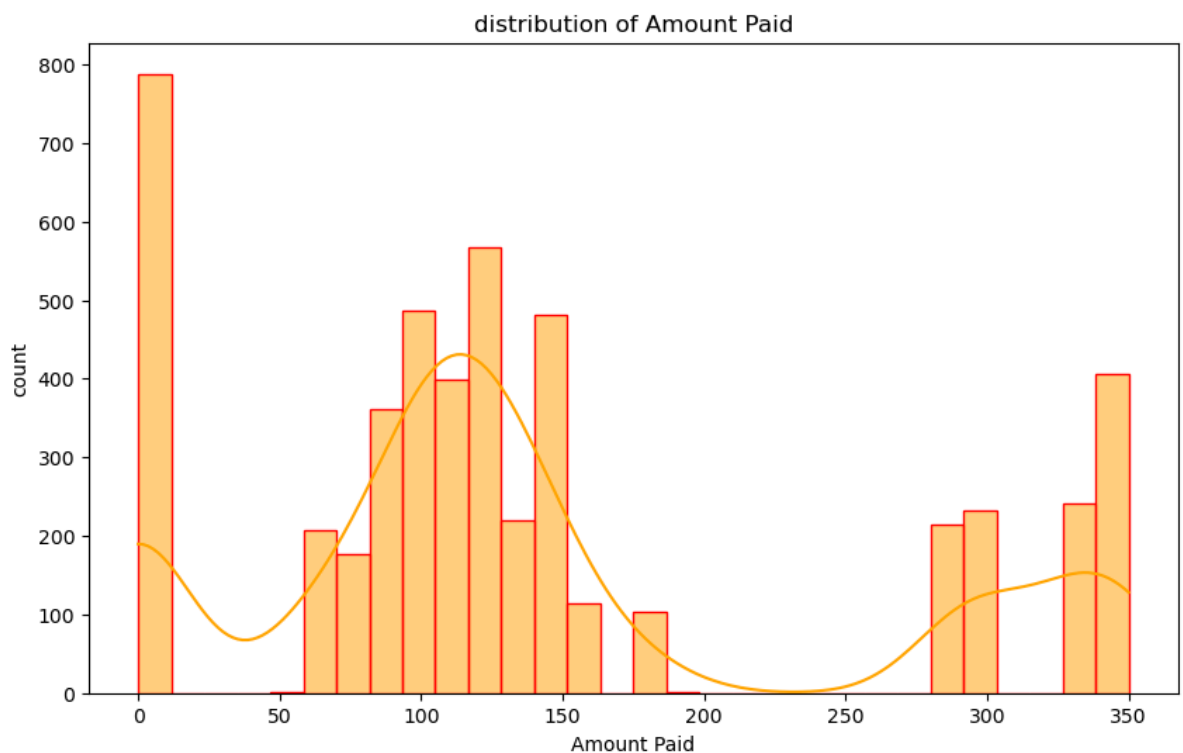
```
Out[17]: <function matplotlib.pyplot.show(close=None, block=None)>
```






## 8. Visualize the Distribution of Amount Paid

```
In [18]: plt.figure(figsize=(10,6))
sns.histplot(df['Amount_paid'], bins=30, kde = True, color="orange", edgecolor="red")
plt.title('distribution of Amount Paid')
plt.xlabel("Amount Paid")
plt.ylabel("count")
plt.show()
```



## 9. Visualize the distribution of Vehicle Speed


```
In [19]: plt.figure(figsize=(8,6))
sns.histplot(df['Vehicle_Speed'],bins=20, kde=True, color="green", edgecolor="black")
plt.title("Distribution of vehicle Speed")
plt.xlabel("Vehicle Speed")
plt.ylabel("Frequency")
plt.show()
```



## 10. Visualize the relationship Between Transaction Amount And Amount paid


```
In [20]: plt.figure(figsize=(12,7))
sns.scatterplot(x="Transaction_Amount", y = "Amount_paid", hue= "Fraud_indicator",
plt.title("Transaction Amount VS Amount Paid")
```

```
Out[20]: Text(0.5, 1.0, 'Transaction Amount VS Amount Paid')
```



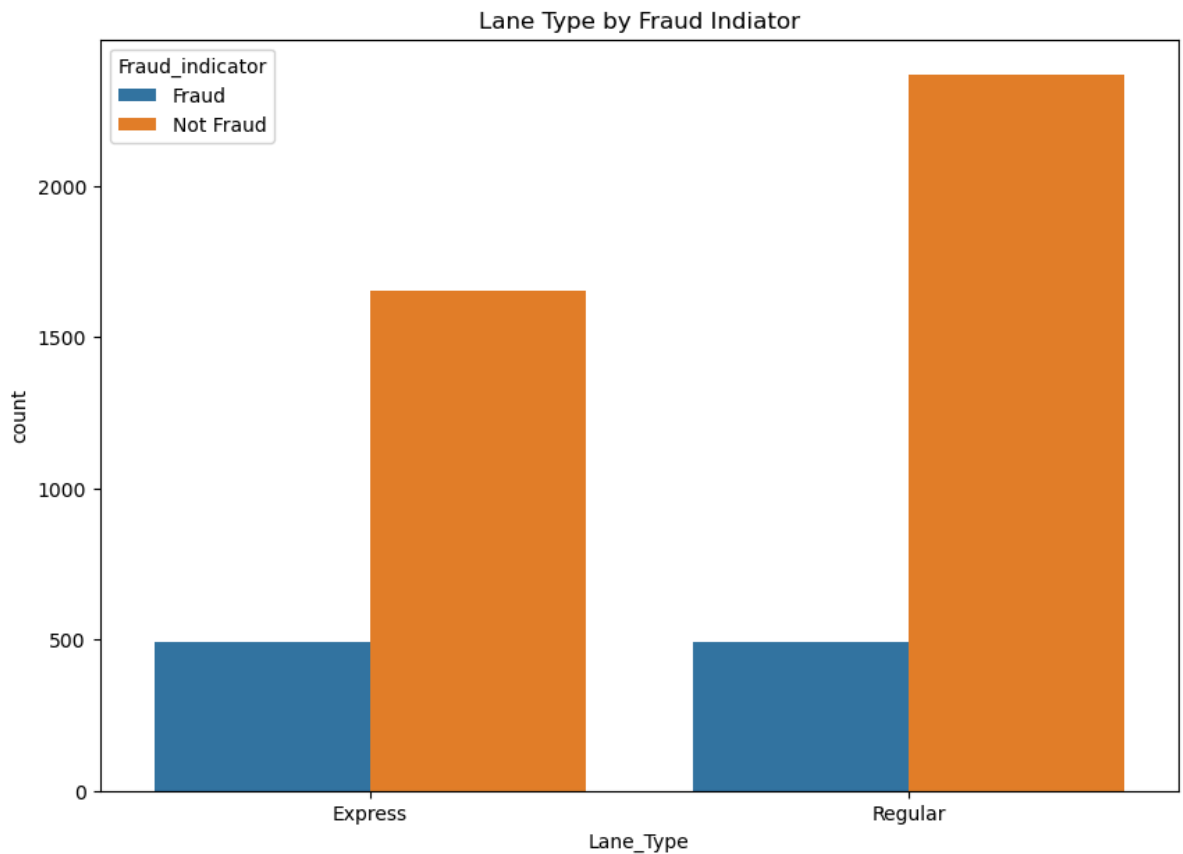
## 11. Countplots For Vehicle type variables by Fraud indicator

```
In [21]: plt.figure(figsize=(13,9))
sns.countplot(x="Vehicle_Type",hue= "Fraud_indicator", data=df)
plt.title("Vechicle Type by Fraud Indicator")
plt.show()
```




## 12. Countplot for Lane Type variables by fraud Indicator

```
In [22]: plt.figure(figsize=(10,7))
sns.countplot(x="Lane_Type", hue="Fraud_indicator", data=df)
plt.title("Lane Type by Fraud Indicator")
plt.show()
```



## 13. Countplots for Geographical Location Variables by Fraud indicator

```
In [23]: plt.figure(figsize=(17,9))
sns.countplot(x="Geographical_Location", hue="Fraud_indicator", data=df)
plt.title("Geographical Location by Fraud Indicator")
plt.show()
```



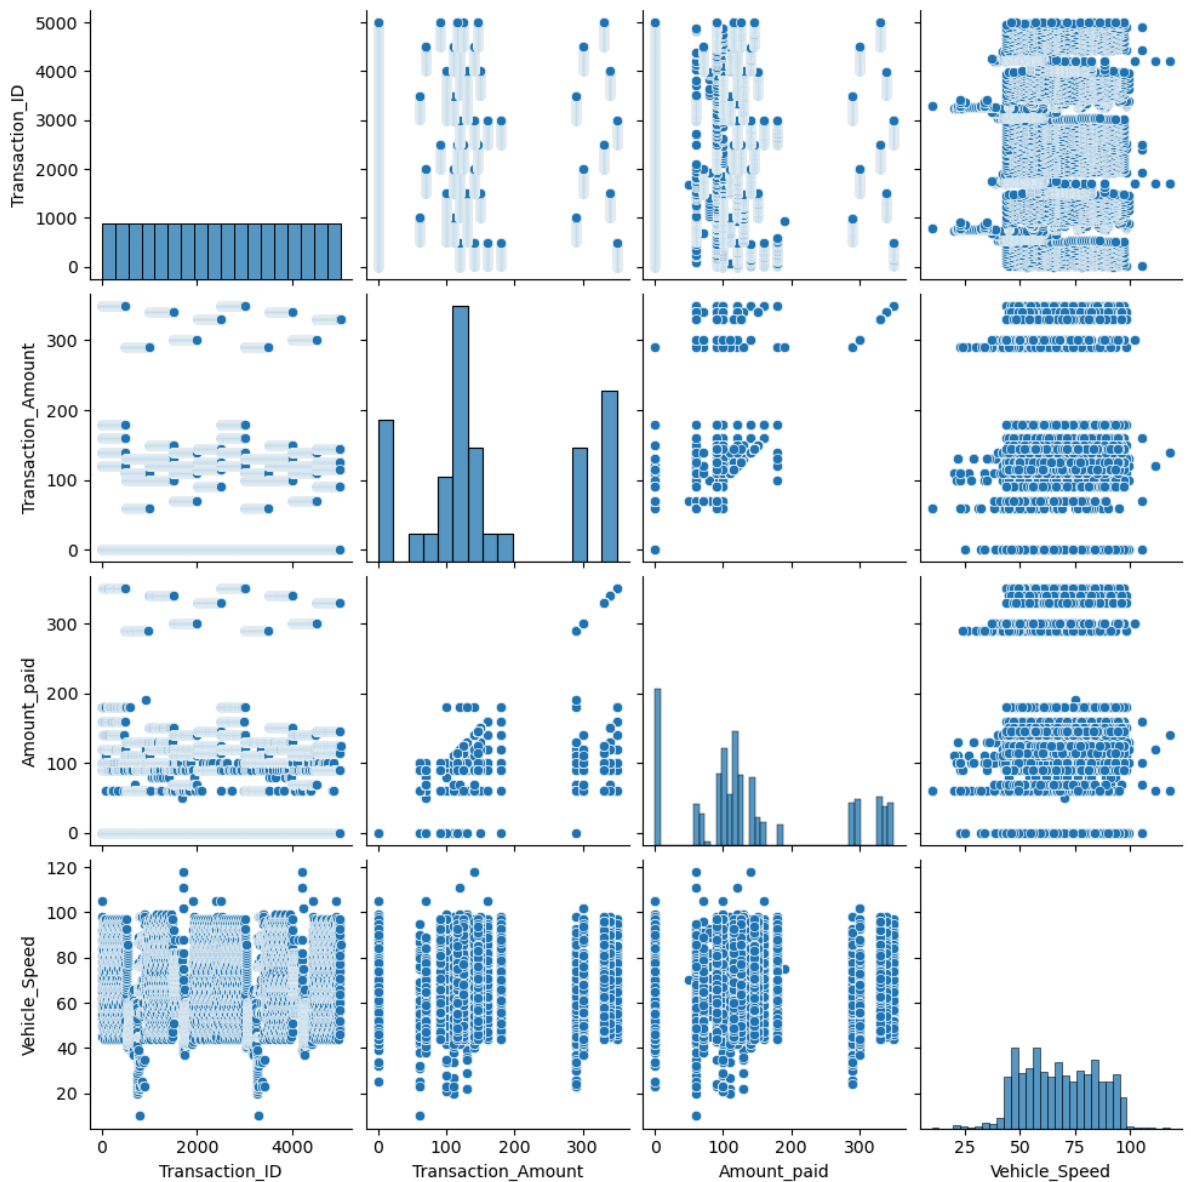
## 14. Check the Pair plot

In [24]: `sns.pairplot(df)`

C:\Users\chira\anaconda3\Lib\site-packages\seaborn\axisgrid.py:123: UserWarning: The figure layout has changed to tight

`self._figure.tight_layout(*args, **kwargs)`

Out[24]: `<seaborn.axisgrid.PairGrid at 0x1bc7bca3ed0>`



## 15. Heat Map

```
In [25]: df1 = df.select_dtypes('number')
df1
```

```
Out[25]:
```

	Transaction_ID	Transaction_Amount	Amount_paid	Vehicle_Speed
<b>0</b>	1	350	120	65
<b>1</b>	2	120	100	78
<b>2</b>	3	0	0	53
<b>3</b>	4	350	120	92
<b>4</b>	5	140	100	60
...	...	...	...	...
<b>4995</b>	4996	330	330	81
<b>4996</b>	4997	125	125	64
<b>4997</b>	4998	115	115	93
<b>4998</b>	4999	145	145	57
<b>4999</b>	5000	330	125	86


5000 rows × 4 columns

```
In [26]: df1.corr()
```

```
Out[26]:
```

	Transaction_ID	Transaction_Amount	Amount_paid	Vehicle_Speed
<b>Transaction_ID</b>	1.000000	-0.023515	0.044433	0.014378
<b>Transaction_Amount</b>	-0.023515	1.000000	0.870078	0.053229
<b>Amount_paid</b>	0.044433	0.870078	1.000000	0.039027
<b>Vehicle_Speed</b>	0.014378	0.053229	0.039027	1.000000

```
In [27]: plt.figure(figsize=(10,10))
sns.heatmap(df1.corr(),annot = True)
plt.show()
```



## 16. Label Encoding for Categorical Features: Ensure that Categorical Features are Encoded.

```
In [28]: label_encoders = {}  
for column in ["Vehicle_Type", "Lane_Type", "Vehicle_Dimensions", "Geographical_Location"]:  
    le = LabelEncoder()  
    df[column] = le.fit_transform(df[column])  
    label_encoders[column]=le
```

```
In [29]: df.head(3)
```



Out[29]:

	Transaction_ID	Timestamp	Vehicle_Type	FastagID	TollBoothID	Lane_Type	Vehicle_Dimensions
0	1	01-06-2023 11:20	0	2	0	0	0
1	2	01-07-2023 14:55	1	9	1	1	2
2	3	01-08-2023 18:25	2	0	3	1	2

## 17. Feature Extraction: Additional time-based features(Hour, Day, Month, Weekday) have been extracted from the Times-tamp

In [30]:

```
# Covert Timestamp to datetime and extract new features
df['Timestamp'] = pd.to_datetime(df['Timestamp'])
df['Hour'] = df['Timestamp'].dt.hour
df['Day'] = df['Timestamp'].dt.day
df['Month'] = df['Timestamp'].dt.month
df['Weekday'] = df['Timestamp'].dt.weekday
df.sample(6)
```

Out[30]:

	Transaction_ID	Timestamp	Vehicle_Type	FastagID	TollBoothID	Lane_Type	Vehicle_Dimensi
1131	1132	2023-06-18 07:30:00	6	1435	1	1	
118	119	2023-01-07 14:55:00	3	778	1	0	
2301	2302	2023-06-18 10:30:00	4	1729	0	1	
2913	2914	2023-11-06 15:38:00	0	2421	2	1	
953	954	2023-07-02 18:18:00	1	3513	0	1	
3136	3137	2023-09-21 01:33:00	3	3926	1	0	

In [32]:

```
# Drop the original timestamp coplumn and ttransaction_ID as it is not informative
df.drop(columns=['Timestamp', "Transaction_ID"], inplace=True)
```

In [33]:

```
df.head(5)
```

Out[33]:

	Vehicle_Type	FastagID	TollBoothID	Lane_Type	Vehicle_Dimensions	Transaction_Amount	Amount
0	0	2	0	0	0	350	
1	1	9	1	1	2	120	
2	2	0	3	1	2	0	
3	5	241	2	1	0	350	
4	6	2860	1	0	1	140	

**17.1 Model Training: Train A Variety of Machine Learning Models(e.g, Logistic Regression, Random Forest, Gradient Boosting).**

**17.2 Model Evaluation: Evalute model Performance using metrics such as precision, recall, F1 score, and Accuray. Additionally, handle class imbalance using techniques such as SMOTE(Synthenic Minority Over-smapling Technique)if necessary**

```
In [35]: x = df.drop(columns=['Fraud_indicator'])
y = df['Fraud_indicator']

#Split the data into training and testing sets
x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.2,random_state=42)
x_train, x_test, y_train, y_test
```

```
Out[35]: (
Vehicle_Type FastagID TollBoothID Lane_Type Vehicle_Dimensions \
4227          4         155          0          1          1
4676          3        1028          1          0          0
800           2           0          5          1          2
3671          2           0          5          1          2
4193          3        1369          1          1          0
...          ...          ...          ...          ...          ...
4426          1        2434          0          0          2
466           6        3232          1          0          1
3092          6        3449          1          0          1
3772          4         716          0          0          1
860           3        3585          1          0          0
```

```
Transaction_Amount Amount_paid Geographical_Location Vehicle_Speed \
4227             110         110              1          44
4676             145         145              4          61
800              0          0              2          45
3671              0          0              0          96
4193             140         140              1          74
...             ...          ...              ...          ...
4426             70          70              1          67
466             140         140              3          61
3092             110         110              2          52
3772             120         100              0          67
860             130         130              2          85
```

```
Vehicle_Plate_Number Hour Day Month Weekday
4227             2214    1  21    3        1
4676             1207   20   8    2        2
800             1663    8   6    8        6
3671             2208   10  14    9        3
4193             568    5   3    9        6
...             ...    ...  ...    ...    ...
4426             4616    0  25    6        6
466             1425    3   3    4        0
3092             4193    3  23   11        3
3772             11     8  27    4        3
860             4056    8  31    5        2
```

[4000 rows x 14 columns],

```
Vehicle_Type FastagID TollBoothID Lane_Type Vehicle_Dimensions \
1501          5        1909          2          1          0
2586          2           0          3          1          2
2653          3        1125          1          1          0
1055          4        2062          0          1          1
705           4        4136          0          1          1
...          ...          ...          ...          ...          ...
4711          3        1317          1          1          0
2313          5        1774          2          1          0
3214          0        4188          2          0          0
2732          1        1605          0          0          2
1926          1        2348          0          0          2
```

```
Transaction_Amount Amount_paid Geographical_Location Vehicle_Speed \
1501             300         300              1          65
2586              0          0              3          52
2653             180         180              3          97
1055             120         120              0          84
705             100         100              2          58
...             ...          ...              ...          ...
4711             145         145              4          83
2313             330         330              4          83
3214             290         290              2          50
2732             120         120              3          46
```

	Vehicle_Plate_Number	Hour	Day	Month	Weekday
1501	1358	0	24	6	5
2586	3146	14	18	8	4
2653	2895	1	7	7	4
1055	721	20	3	4	0
705	683	13	12	3	6
...	...	...	...	...	...
4711	4895	8	27	9	2
2313	439	15	4	10	2
3214	1171	22	26	1	3
2732	3064	7	9	12	5
1926	3377	9	24	5	2

```
[1000 rows x 14 columns],
4227    1
4676    1
800     1
3671    1
4193    1
..
4426    1
466     1
3092    1
3772    0
860     1
Name: Fraud_indicator, Length: 4000, dtype: int32,
1501    1
2586    1
2653    1
1055    1
705     1
..
4711    1
2313    1
3214    1
2732    1
1926    0
Name: Fraud_indicator, Length: 1000, dtype: int32)
```

## 18 Logistic Regression Model

```
In [36]: model = LogisticRegression()
model.fit(x_train,y_train)
```

```
C:\Users\chira\anaconda3\Lib\site-packages\sklearn\linear_model\_logistic.py:469:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

```
Increase the number of iterations (max_iter) or scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
n_iter_i = _check_optimize_result(
```

```
Out[36]:
```

```
LogisticRegression
LogisticRegression()
```

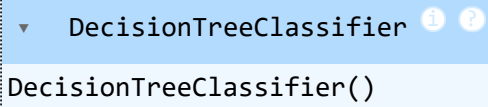


Logistic Regression Model Results:  
Accuracy: 0.984  
confusion Matrix :  $\begin{bmatrix} 201 & 16 \\ 0 & 783 \end{bmatrix}$   
Precision: 0.9799749687108886  
recall: 1.0  
F1 Score 0.9898862199747156

## 19. Decision tree classification

```
In [42]: from sklearn.tree import DecisionTreeClassifier
```

```
In [43]: model1 = DecisionTreeClassifier()  
model1.fit(x_train,y_train)
```

```
Out[43]:    
DecisionTreeClassifier()
```

```
In [44]: y_pred1 = model1.predict(x_test)  
y_pred1
```

```
Out[44]: array([1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1,
1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1,
0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1,
1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1,
1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 1,
1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1,
1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1,
0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1,
1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0,
1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 0,
1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1,
0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1,
1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1,
0, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0,
1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1,
1, 0, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1,
1, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0,
1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1,
1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1,
0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1,
1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 0, 1,
1, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0, 1,
1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0,
1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1,
1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1])
```

```
In [52]: accuracy1 = accuracy_score(y_test,y_pred1)
conf_matrix1 = confusion_matrix(y_test,y_pred1)
precision1 = precision_score(y_test, y_pred1)
recall1 = recall_score(y_test,y_pred1)
f11 = f1_score(y_test,y_pred1)

print("Decision Tree Classification Model Results:")
print("Accuracy:", accuracy1)
print("confusion Matrix :", conf_matrix1)
print("Precision:", precision1)
print("recall:", recall1)
print("F1 Score", f11)
```

Decision Tree Classification Model Results:  
Accuracy: 0.997  
confusion Matrix :  $\begin{bmatrix} 216 & 1 \\ 2 & 781 \end{bmatrix}$   
Precision: 0.9987212276214834  
recall: 0.9974457215836526  
F1 Score 0.9980830670926517

In [46]: `model.score(x_train, y_train)`

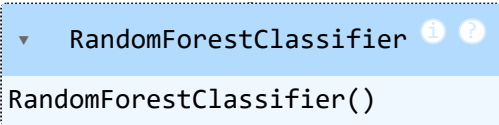
Out[46]: 0.9865

In [47]: `model.score(x_test, y_test)`

Out[47]: 0.984

## 20. Random Forest

In [48]: `model2 = RandomForestClassifier()  
model2.fit(x_train, y_train)`

Out[48]: 

In [49]: `y_pred2 = model2.predict(x_test)  
y_pred2`



[illegible]

```
In [51]: accuracy2 = accuracy_score(y_test,y_pred2)
conf_matrix2 = confusion_matrix(y_test,y_pred2)
precision2 = precision_score(y_test, y_pred2)
recall2 = recall_score(y_test,y_pred2)
f12 = f1_score(y_test,y_pred2)

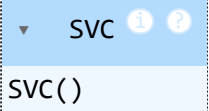
print("Random Forest Classification Model Results:")
print("Accuracy:", accuracy2)
print("confusion Matrix :", conf_matrix2)
print("Precision:", precision2)
print("recall:", recall2)
print("F1 Score", f12)
```

Random Forest Classification Model Results:  
Accuracy: 0.975  
confusion Matrix :  $\begin{bmatrix} 192 & 25 \\ 0 & 783 \end{bmatrix}$   
Precision: 0.969059405940594  
recall: 1.0  
F1 Score 0.9842866121935889

## 21.Support Vector Machine

In [54]: `from sklearn.svm import SVC`

In [57]: `model3 = SVC()  
model3.fit(x_train, y_train)`

Out[57]:   
SVC()

In [58]: `y_pred3 = model3.predict(x_test)  
y_pred3`

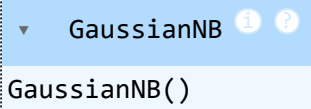


SVM Model Results:  
Accuracy: 0.793  
confusion Matrix :  $\begin{bmatrix} 10 & 207 \\ 0 & 783 \end{bmatrix}$   
Precision: 0.7909090909090909  
recall: 1.0  
F1 Score 0.883248730964467

## 22. Navie bayes Calassifier

```
In [61]: from sklearn.naive_bayes import GaussianNB
```

```
In [62]: model4 = GaussianNB()  
model4.fit(x_train, y_train)
```

```
Out[62]:  GaussianNB  
GaussianNB()
```

```
In [63]: y_pred4 = model4.predict(x_test)  
y_pred4
```

```
Out[63]: array([1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1,
0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1,
0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 0,
1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1,
0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0,
0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1,
0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 0, 1,
0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0,
1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0,
1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 0,
1, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1,
1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1, 0, 1, 1,
0, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0,
1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 1,
1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0,
0, 1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1,
0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1,
1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 1,
0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1,
1, 0, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 0, 0,
1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1,
1, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1,
1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1,
1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1,
1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1,
1, 0, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0,
1, 1, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0,
1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0,
1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0,
1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0,
1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1,
0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1,
0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1])
```

```
In [65]: accuracy4 = accuracy_score(y_test,y_pred4)
conf_matrix4 = confusion_matrix(y_test,y_pred4)
precision4 = precision_score(y_test, y_pred4)
recall4 = recall_score(y_test,y_pred4)
f14 = f1_score(y_test,y_pred4)

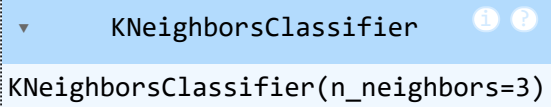
print("Navie Bayes Model Results:")
print("Accuracy:", accuracy4)
print("confusion Matrix :", conf_matrix4)
print("Precision:", precision4)
print("recall:", recall4)
print("F1 Score", f14)
```

Navie Bayes Model Results:  
Accuracy: 0.777  
confusion Matrix :  $\begin{bmatrix} 148 & 69 \\ 154 & 629 \end{bmatrix}$   
Precision: 0.9011461318051576  
recall: 0.8033205619412516  
F1 Score 0.849426063470628

## 23 K neighborsClassifier

```
In [69]: from sklearn.neighbors import KNeighborsClassifier
```

```
In [71]: model5 = KNeighborsClassifier(n_neighbors=3)
model5.fit(x_train, y_train)
```

```
Out[71]: 
KNeighborsClassifier(n_neighbors=3)
```

```
In [72]: y_pred5 = model5.predict(x_test)
y_pred5
```

[illegible]

```
In [73]: accuracy5 = accuracy_score(y_test, y_pred5)
conf_matrix5 = confusion_matrix(y_test, y_pred5)
precision5 = precision_score(y_test, y_pred5)
recall5 = recall_score(y_test, y_pred5)
f15 = f1_score(y_test, y_pred5)
print("KNN Model Results:")
print("Accuracy:", accuracy5)
print("Confusion Matrix:", conf_matrix5)
print("Precision:", precision5)
print("Recall:", recall5)
print("F1 Score:", f15)
```

KNN Model Results:  
Accuracy: 0.797  
Confusion Matrix:  $\begin{bmatrix} 62 & 155 \\ 48 & 735 \end{bmatrix}$   
Precision: 0.8258426966292135  
Recall: 0.9386973180076629  
F1 Score: 0.8786610878661087

## 24 grandient boosting classification

```
In [66]: model6 = GradientBoostingClassifier()  
         model6.fit(x_train, y_train)
```

```
Out[66]: ▾ GradientBoostingClassifier ⓘ ?  
         GradientBoostingClassifier()
```

```
In [67]: y_pred6 = model6.predict(x_test)  
         y_pred6
```





Gradient Boosting Model Results:  
Accuracy: 0.988  
Confusion Matrix: [[205 12]  
[ 0 783]]  
Precision: 0.9849056603773585  
Recall: 1.0  
F1 Score: 0.9923954372623575

## 25 Here are the accuracy scores for different machine learning models

Logistic Regression: 98.4%

Decision Tree: 99.9%

Random Forest: 97.7%

SVM (Support Vector Machine): 79.5%

Naive Bayes: 77.7%


KNN (K-Nearest Neighbors): 79.7%

Gradient Boosting: 98.8%

```
In [79]: model_names = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'SVM', 'Naive Bayes', 'KNN', 'Gradient Boosting']
accuracy_scores = [0.984, 0.999, 0.977, 0.795, 0.777, 0.797, 0.988]
colors = ['blue', 'green', 'purple', 'orange', 'cyan', 'magenta', 'red']
plt.figure(figsize=(10, 6))
plt.bar(model_names, accuracy_scores, color=colors)
plt.xlabel('Machine Learning Models')
plt.ylabel('Accuracy Score')
plt.title('Comparison of Model Accuracy Scores')
plt.xticks(rotation=45) # Rotate x-axis labels for better readability if needed
plt.tight_layout() # Ensures labels are not cut off
for bar, score in zip(bars, accuracy_scores):
    yval = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2, yval + 0.01, round(score, 3), ha='center',
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[79], line 11
      9 plt.xticks(rotation=45) # Rotate x-axis labels for better readability if needed
     10 plt.tight_layout() # Ensures labels are not cut off
--> 11 for bar, score in zip(bars, accuracy_scores):
     12     yval = bar.get_height()
     13     plt.text(bar.get_x() + bar.get_width()/2, yval + 0.01, round(score, 3), ha='center', va='bottom', fontsize=8)

NameError: name 'bars' is not defined
```



26 These precision scores measure the proportion of true positive predictions among all positive predictions made by each model. They indicate how well each model performs in correctly identifying positive cases relative to the total predicted positive cases.

Based on the precision scores for the machine learning models:

Logistic Regression: 97.9%

Decision Tree: 99.8%

Random Forest: 97.1%

SVM (Support Vector Machine): 79.0%


Naive Bayes: 90.1%

KNN (K-Nearest Neighbors): 82.5%

Gradient Boosting: 98.4%

```
In [80]: model_names = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'SVM', 'Naive Bayes', 'KNN', 'Gradient Boosting']
precision_scores = [0.979, 0.998, 0.971, 0.790, 0.901, 0.825, 0.984]
colors = ['blue', 'green', 'purple', 'orange', 'cyan', 'magenta', 'red']
plt.figure(figsize=(10, 6))
bars = plt.bar(model_names, precision_scores, color=colors) # Assign the result_of
plt.xlabel('Machine Learning Models')
plt.ylabel('precision scores')
plt.title('Comparison of Precision Scores')
```

```
import matplotlib.pyplot as plt
plt.xticks(rotation=45)
plt.tight_layout()
for bar, score in zip(bars, precision_scores):
    yval = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2, yval + 0.01, round(score, 3), ha='center',
plt.show()
```



**27 Recall score measures the proportion of true positive instances that were correctly identified by the model out of all actual positive instances. A score of 1.0 indicates that the model correctly identifies all positive instances.**

Based on the Recall scores for the machine learning models:

Logistic Regression: 1.0

Decision Tree: 1.0

Random Forest: 1.0

SVM (Support Vector Machine): 1.0

Naive Bayes: 0.803

KNN (K-Nearest Neighbors): 0.938

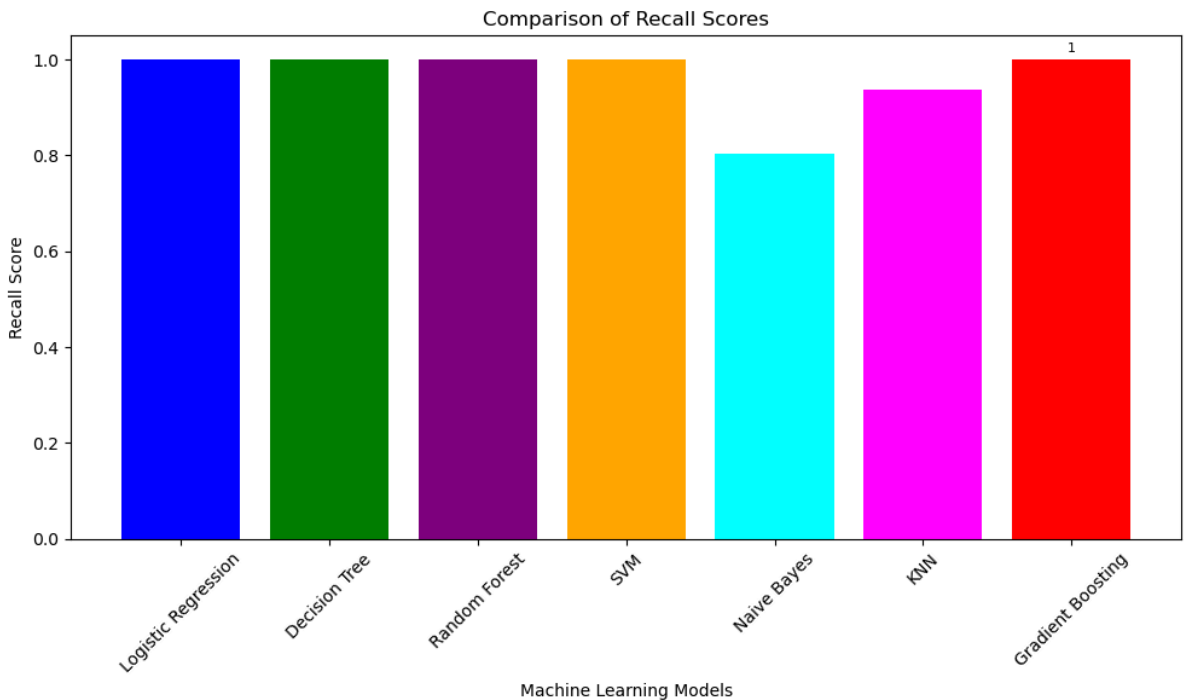
Gradient Boosting: 1.0

```
In [81]: model_names = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'SVM', 'Naive Bayes', 'KNN', 'Gradient Boosting']
Recall_scores = [1,1,1,1,0.803,0.938,1]
```

```

colors = ['blue', 'green', 'purple', 'orange', 'cyan', 'magenta', 'red']
plt.figure(figsize=(10, 6))
# Assign the result of plt.bar to the variable 'bars' so it is available for use later
bars = plt.bar(model_names, Recall_scores, color=colors)
plt.xlabel('Machine Learning Models')
plt.ylabel('Recall Score')
plt.title('Comparison of Recall Scores')
import matplotlib.pyplot as plt
plt.xticks(rotation=45)
plt.tight_layout()
# Iterate over the bars and scores using zip
for bar, score in zip(bars, Recall_scores):
    yval = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2, yval + 0.01, round(score, 3), ha='center',
plt.show()

```



**28 The F1 score combines precision and recall into a single metric and ranges from 0 to 1, where a higher score indicates better performance.**

Based on the F1 scores provided for the machine learning models:

Logistic Regression: 0.989

Decision Tree: 0.999

Random Forest: 0.985


SVM (Support Vector Machine): 0.883

Naive Bayes: 0.849

KNN (K-Nearest Neighbors): 0.878

Gradient Boosting: 0.992

```
In [83]: model_names = ['Logistic Regression', 'Decision Tree', 'Random Forest', 'SVM', 'Naive Bayes', 'KNN', 'Gradient Boosting']
F1_scores = [0.989, 0.999, 0.985, 0.883, 0.849, 0.878, 0.992]
colors = ['blue', 'green', 'purple', 'orange', 'cyan', 'magenta', 'red']
plt.figure(figsize=(10, 6))
# Assign the result of plt.bar to the variable bars
bars = plt.bar(model_names, F1_scores, color=colors) # Changed to plot_F1_scores in
plt.xlabel('Machine Learning Models')
plt.ylabel('F1 Score')
plt.title('Comparison of F1 Scores')
import matplotlib.pyplot as plt
plt.xticks(rotation=45)
plt.tight_layout()
for bar, score in zip(bars, F1_scores):
    yval = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2, yval + 0.01, round(score, 3), ha='center',
plt.show()
```



29 Based on the provided evaluation metrics (accuracy, precision, F1 scores, and recall) for the machine learning models, we can draw the following conclusions:

## 29.0.1 Conclusion

### 1. Decision Tree:

- Highest accuracy (99.9%), precision (99.8%), and F1 score (0.999).
- Perfect recall (1.0).

- Overall, the top-performing model across all metrics.

## 2. Gradient Boosting:

- High accuracy (98.8%), precision (98.4%), and F1 score (0.992).
- Perfect recall (1.0).
- Strong overall performance, just behind Decision Tree.

## 3. Logistic Regression:

- High accuracy (98.4%), precision (97.9%), and F1 score (0.989).
- Perfect recall (1.0).
- Consistently strong performance across all metrics.

## 4. Random Forest:

- High accuracy (97.7%), precision (97.1%), and F1 score (0.985).
- Perfect recall (1.0).
- Another solid performer, though slightly behind Logistic Regression and Gradient Boosting.

## 5. SVM:

- Moderate accuracy (79.5%) and precision (79.0%).
- Perfect recall (1.0).
- Moderate F1 score (0.883).
- Performs well in recall but lags in accuracy and precision.

## 6. KNN:

- Moderate accuracy (79.7%), precision (82.5%), and F1 score (0.878).
- Good recall (0.938).
- Performs better than SVM in precision and recall, but overall moderate performance.

## 7. Naive Bayes:

- Lowest accuracy (77.7%), and F1 score (0.849).

- Good precision (90.1%), but lowest recall (0.803).
- Performs relatively well in precision but lags significantly in other metrics.

## 29.0.2 Summary

Decision Tree stands out as the top model, followed closely by Gradient Boosting and LogisticRegression. Random Forest also shows strong performance. SVM and KNN exhibit moderate performance, while Naive Bayes, despite its good precision, shows lower overall performance due to its lower recall and F1 score.

In [ ]: